

▼ Predicting Chronic Kidney Disease based on health records

Given 24 health related attributes taken in 2-month period of 400 patients, using the information of the 158 patients with complete records to predict the outcome (i.e. whether one has chronic kidney disease) of the remaining 242 patients (with missing values in their records).

▼ Load Modules and helper functions

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import roc_curve, auc, confusion_matrix, classification_report, accuracy_score
from sklearn.ensemble import RandomForestClassifier
import warnings
warnings.filterwarnings('ignore')

# from subprocess import check_output
# print(check_output(["ls", "../input"]).decode("utf8"))

%matplotlib inline

def auc_scorer(clf, X, y, model): # Helper function to plot the ROC curve
    if model=='RF':
        fpr, tpr, _ = roc_curve(y, clf.predict_proba(X)[:,:1])
    elif model=='SVM':
        fpr, tpr, _ = roc_curve(y, clf.decision_function(X))
    roc_auc = auc(fpr, tpr)

    plt.figure()    # Plot the ROC curve
```

```
plt.plot(fpr, tpr, label='ROC curve from '+model+' model (area = %0.3f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc="lower right")
plt.show()

return fpr,tpr,roc_auc
```

▼ Load files

```
df = pd.read_csv("C:/Users/Sinegalatha/Desktop/2nd year online class/nalaiya thiran/dataset/kidney_disease.csv")
```

```
df.head()
```

	id	age	bp	sg	al	su	rbc	pc	pcc	ba	...	pcv	wc	rc	htn	dm	cad	appet	pe	ane	cla
0	0	48.0	80.0	1.020	1.0	0.0	NaN	normal	notpresent	notpresent	...	44	7800	5.2	yes	yes	no	good	no	no	
1	1	7.0	50.0	1.020	4.0	0.0	NaN	normal	notpresent	notpresent	...	38	6000	NaN	no	no	no	good	no	no	
2	2	62.0	80.0	1.010	2.0	3.0	normal	normal	notpresent	notpresent	...	31	7500	NaN	no	yes	no	poor	no	yes	
3	3	48.0	70.0	1.005	4.0	0.0	normal	abnormal	present	notpresent	...	32	6700	3.9	yes	no	no	poor	yes	yes	
4	4	51.0	80.0	1.010	2.0	0.0	normal	normal	notpresent	notpresent	...	35	7300	4.6	no	no	no	good	no	no	

5 rows × 26 columns



```
df['wc']
```

```

0      7800
1      6000
2      7500
3      6700
4      7300
...
395    6700
396    7800
397    6600
398    7200
399    6800

```

Name: wc, Length: 400, dtype: object

df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 400 entries, 0 to 399

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	id	400 non-null	int64
1	age	391 non-null	float64
2	bp	388 non-null	float64
3	sg	353 non-null	float64
4	al	354 non-null	float64
5	su	351 non-null	float64
6	rbc	248 non-null	object
7	pc	335 non-null	object
8	pcc	396 non-null	object
9	ba	396 non-null	object
10	bgr	356 non-null	float64
11	bu	381 non-null	float64
12	sc	383 non-null	float64
13	sod	313 non-null	float64
14	pot	312 non-null	float64
15	hemo	348 non-null	float64
16	pcv	330 non-null	object
17	wc	295 non-null	object
18	rc	270 non-null	object

```

19  htn          398 non-null  object
20  dm           398 non-null  object
21  cad          398 non-null  object
22  appet        399 non-null  object
23  pe           399 non-null  object
24  ane          399 non-null  object
25  classification 400 non-null  object
dtypes: float64(11), int64(1), object(14)
memory usage: 81.4+ KB

```

```
df.describe()
```

	id	age	bp	sg	al	su	bgr	bu	sc	sod	
count	400.000000	391.000000	388.000000	353.000000	354.000000	351.000000	356.000000	381.000000	383.000000	313.000000	312.000000
mean	199.500000	51.483376	76.469072	1.017408	1.016949	0.450142	148.036517	57.425722	3.072454	137.528754	4.625000
std	115.614301	17.169714	13.683637	0.005717	1.352679	1.099191	79.281714	50.503006	5.741126	10.408752	3.150000
min	0.000000	2.000000	50.000000	1.005000	0.000000	0.000000	22.000000	1.500000	0.400000	4.500000	2.500000
25%	99.750000	42.000000	70.000000	1.010000	0.000000	0.000000	99.000000	27.000000	0.900000	135.000000	3.800000
50%	199.500000	55.000000	80.000000	1.020000	0.000000	0.000000	121.000000	42.000000	1.300000	138.000000	4.400000
75%	299.250000	64.500000	80.000000	1.020000	2.000000	0.000000	163.000000	66.000000	2.800000	142.000000	4.900000
max	399.000000	90.000000	180.000000	1.025000	5.000000	5.000000	490.000000	391.000000	76.000000	163.000000	47.000000

```
df[df.duplicated()]
```

```

   id  age  bp  sg  al  su  rbc  pc  pcc  ba  ...  pcv  wc  rc  htn  dm  cad  appet  pe  ane  classification
0 rows × 26 columns

```

```
df.isna().sum()
```

id	0
age	9
bp	12
sg	47
al	46
su	49
rbc	152
pc	65
pcc	4
ba	4
bgr	44
bu	19
sc	17
sod	87
pot	88
hemo	52
pcv	70
wc	105
rc	130
htn	2
dm	2
cad	2
appet	1
pe	1
ane	1
classification	0
dtype:	int64

```
df2 = df.dropna(axis=0)
```

▼ Cleaning and preprocessing of data for training a classifier

```
# Map text to 1/0 and do some cleaning
df[['htn','dm','cad','pe','ane']] = df[['htn','dm','cad','pe','ane']].replace(to_replace={'yes':1,'no':0})
df[['rbc','pc']] = df[['rbc','pc']].replace(to_replace={'abnormal':1,'normal':0})
df[['pcc','ba']] = df[['pcc','ba']].replace(to_replace={'present':1,'notpresent':0})
```

```
df[['appet']] = df[['appet']].replace(to_replace={'good':1,'poor':0,'no':np.nan})
df['classification'] = df['classification'].replace(to_replace={'ckd':1.0,'ckd\t':1.0,'notckd':0.0,'no':0.0})
df.rename(columns={'classification':'class'},inplace=True)
```

```
# Further cleaning
df['pe'] = df['pe'].replace(to_replace='good',value=0) # Not having pedal edema is good
df['appet'] = df['appet'].replace(to_replace='no',value=0)
df['cad'] = df['cad'].replace(to_replace='\tno',value=0)
df['dm'] = df['dm'].replace(to_replace={'\tno':0,'\tyes':1,' yes':1, '':np.nan})
df.drop('id',axis=1,inplace=True)
```

```
df.head()
```

	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	...	pcv	wc	rc	htn	dm	cad	appet	pe	ane	class
0	48.0	80.0	1.020	1.0	0.0	NaN	0.0	0.0	0.0	121.0	...	44	7800	5.2	1.0	1.0	0.0	1.0	0.0	0.0	1.0
1	7.0	50.0	1.020	4.0	0.0	NaN	0.0	0.0	0.0	NaN	...	38	6000	NaN	0.0	0.0	0.0	1.0	0.0	0.0	1.0
2	62.0	80.0	1.010	2.0	3.0	0.0	0.0	0.0	0.0	423.0	...	31	7500	NaN	0.0	1.0	0.0	0.0	0.0	1.0	1.0
3	48.0	70.0	1.005	4.0	0.0	0.0	1.0	1.0	0.0	117.0	...	32	6700	3.9	1.0	0.0	0.0	0.0	1.0	1.0	1.0
4	51.0	80.0	1.010	2.0	0.0	0.0	0.0	0.0	0.0	106.0	...	35	7300	4.6	0.0	0.0	0.0	1.0	0.0	0.0	1.0

5 rows × 25 columns

► Check the portion of rows with NaN

- Now the data is cleaned with improper values labelled NaN. Let's see how many NaNs are there.
- Drop all the rows with NaN values, and build a model out of this dataset (i.e. df2)

[] ↳ 1 cell hidden

- ▶ Examine correlations between different features

[] ↳ 1 cell hidden

- ▶ Split the set for training models further into a (sub-)training set and testing set.

[] ↳ 4 cells hidden

- ▶ Choosing parameters with GridSearchCV with 10-fold cross validations.

(Suggestion for next time: try using Bayesian model selection method)

[] ↳ 1 cell hidden

- ▶ Examine feature importance

Since I pruned the forest ($max_depth=2$) and decrease the number of trees ($n_estimators=8$), not all features are used.

[] ↳ 2 cells hidden

- ▶ Next, I examine the rest of the dataset (with missing values across the rows)

Are there correlations between occurrence of missing values in a row? The plot suggests, seems no.

[] ↳ 1 cell hidden

- ▶ Make predictions with the best model selected above

I filled in all NaN with 0 and pass it to the trained classifier. The results are as follows:

- True positive = 180
- True negative = 35
- False positive = 0
- False negative = 27

-
- Accuracy = 88.8%
 - ROC AUC = 99.2%

[] ↳ 2 cells hidden

▼ RandomForest Regressor

```
from sklearn.ensemble import RandomForestRegressor
reg=RandomForestRegressor()
reg.fit(X_train,y_train)
```

```
▼ RandomForestRegressor
RandomForestRegressor()
```

```
y_pred=reg.predict(X_test)
```

```
pickle. dump(reg, open('randomreg_chronic', 'wb'))
```

```
y_pred
```

```
array([0.6 , 0.8 , 1.  , 0.78, 0.79, 0.29, 0.74, 1.  , 1.  , 1.  , 0.29,
       1.  , 1.  , 0.29, 0.09, 0.88, 0.29, 0.09, 0.77, 0.29, 0.29, 0.69,
```



```
0.66, 0.29, 0.99, 0.88, 1. , 0.89, 0.98, 0.89, 0.29, 1. , 0.94,
1. , 0.09, 0.07, 0.89, 1. , 1. , 0.07, 0.73, 1. , 0.29, 0.99,
0.29, 0.01, 0.94, 0.8 , 1. , 0.29, 0.29, 0.89, 0.89, 0.8 , 0.89,
0.09, 0.62, 0.99, 0.8 , 1. , 0. , 0. , 0.8 , 1. , 0.69, 0.89,
0.29, 0.81, 0.29, 0.29, 0.29, 0.69, 0.99, 0.29, 0.87, 0.98, 0.09,
0.28, 0.69, 0.89, 0.89, 0.29, 0.09, 0.8 , 1. , 0.22, 1. , 0.29,
0.2 , 0.29, 0.89, 0.09, 0.29, 0.27, 1. , 0.8 , 0.09, 1. , 0.05,
0.8 , 0.09, 0.8 , 0.09, 0.89, 0.73, 0.29, 0.73, 0.29, 0.29, 0.27,
0.69, 0.09, 0.29, 0.29, 0.05, 0.29, 1. , 0.88, 1. , 0.09, 0.89,
0.29, 0.89, 1. , 0.69, 0.29, 0.69, 0.76, 0.28, 0.09, 1. , 1. ,
1. , 0.99, 0.29, 0.71, 0.28, 0.29, 0.01, 0.09, 0.29, 0.05, 0.98,
0.29, 0.99, 0.87, 0.82, 0.27, 0.29, 0.93, 1. , 0.99, 0.89, 0.09,
0.29, 1. , 0.69, 0.8 , 0.8 , 0.29, 1. , 0.09, 0.89, 0.8 , 1. ,
0.29, 0.28, 0.89, 0.29, 0.29, 0.29, 1. , 0.25, 0.82, 0.23, 0.22,
0.09, 0.09, 0. , 0.8 , 0.09, 0.89, 0.09, 0.29, 0.09, 0.88, 0.29,
0.04, 0.29, 0.28, 0.29, 0.99, 0.29, 0.69, 0.99, 0. , 1. , 0.29,
0.97, 0.29, 0.98, 0.69, 0.88, 0.91, 0.95, 0.29, 0.69, 0. , 0.09,
0.02, 0. , 0. , 0. , 0.02, 0. , 0. , 0. , 0.01, 0. , 0. ,
0. , 0.02, 0. , 0.01, 0. , 0.02, 0. , 0.09, 0. , 0.09, 0.03,
0.05, 0. , 0. , 0.01, 0. , 0. , 0. , 0. , 0.07, 0. , 0.01])
```

```
l_pred=list(y_pred)
```

```
l_test=list(y_test)
```

```
d={'prob':l_pred,'out':y_test}
```

```
df_i=pd.DataFrame(d)
```

```
df_i.head()
```

	prob	out
0	0.60	1.0
1	0.80	1.0
2	1.00	1.0

```
df_i.to_csv('C:/Users/Sinegalatha/Desktop/2nd year online class/nalaiya thiran/output/file1.csv')
```

5	0.79	1.0
---	------	-----

► Summary of Results

With proper tuning of parameters using cross-validation in the training set, the Random Forest Classifier achieves an accuracy of 88.8% and an ROC AUC of 99.2%. Lesson learnt: It happens that some pruning helps improve the performance of RF a lot.

[] ↳ 3 cells hidden